

Determining the alternatives for scalar implicature

Benjamin Peloquin

bpeloqui@stanford.edu

Department of Psychology

Stanford University

Michael C. Frank

mcf Frank@stanford.edu

Department of Psychology

Stanford University

Abstract

Successful communication regularly requires listeners to make pragmatic inferences - enrichments beyond the literal meaning of a speaker's utterance. For example, when interpreting a sentence such as "Alice ate some of the cookies," listeners routinely infer that Alice did not eat all of them. A Gricean account of this phenomena assumes the presence of alternatives (like "all of the cookies") with varying degrees of informativity, but it remains an open question precisely what these alternatives are. To address this question, we collect empirical measurements of speaker and listener judgments about varying sets of alternatives across a range of scales and use these as inputs to a computational model of pragmatic inference. This approach allows us to test hypotheses about how well different sets of alternatives predict implicature performance. Our findings suggest that comprehenders likely consider a broader set of alternatives beyond those logically entailed by the initial message.

Keywords: pragmatics; scalar implicature; bayesian modeling

Introduction

How much of what we mean comes from the words that go unsaid? As listeners, our ability to make precise inferences about a speaker's intended meaning in context is indispensable to successful communication. For example, listeners commonly enrich the meaning of the scalar item *some* to *some but not all* in sentences like "Alice ate some of the cookies" (Grice, 1975; Levinson, 2000). These inferences, called *scalar implicatures*, have been an important test case for understanding pragmatic inferences more generally. A Gricean account of this phenomenon assumes listeners reason about a speaker's intended meaning by incorporating knowledge about A) alternative scalar items a speaker could have used (such as *all*) and B) the relative informativity of using such alternatives (Grice, 1975). According to this account, a listener will infer that the speaker must have intended that Alice did not eat all the cookies because it would have been underinformative to use the descriptor *some* when the alternative *all* could have been used just as easily.

But what are the alternatives that should be considered in the implicature computation more generally? Under classic accounts, listeners consider only those words whose meanings entail the actual message (Horn, 1972), and these alternatives enter into conventionalized or semi-conventionalized scales (Levinson, 2000). For example, because *all* entails *some*, and hence is a "stronger" meaning, *all* should be considered as an alternative to *some* in implicatures. Similar scales exist for non-quantifier scales, e.g. *loved* entails *liked* (and hence "I liked the movie" implicates that I didn't love it).

Recent empirical evidence has called into question whether entailment scales are all that is necessary for understanding

scalar implicature. For example, Degen & Tanenhaus (2015) demonstrated that the scalar item *some* was judged less appropriate when exact numbers were seen as viable alternatives. And in a different paradigm, van Tiel (2014) found converging evidence that *some* was judged to be atypical for small quantities. These data provide indirect evidence about a broader set of alternatives. Since *some* is logically true of sets with one or two members, these authors argued that the presence of salient alternatives (the words *one* and *two*, for example) reduced the felicity of *some* via a pragmatic inference.

By formalizing pragmatic reasoning, computational models can help provide more direct evidence about the role that alternatives play. The "rational speech act" model (RSA) is one recent framework for understanding inferences about meaning in context (Frank & Goodman, 2012; Goodman & Stuhlmuller, 2013). RSA models frame language understanding as a special case of social cognition, in which listeners and speakers reason recursively about one another's goals. In the case of scalar implicature, a listener makes a probabilistic inference about what the speaker's most likely communicative goal was, given that she picked the quantifier *some* rather than the stronger quantifier *all*. In turn, the speaker reasons about what message would best convey her intended meaning to the listener, given that he is reasoning in this way. This recursion is grounded in a "literal" listener who reasons only according to the basic truth-functional semantics of the language.

Franke (2014) used an RSA-style model to assess what alternatives a speaker would need to consider in order to produce the typicality/felicity ratings reported by Degen & Tanenhaus (2015) and van Tiel (2014) for the scale *some/all*. In order to do this, Franke (2014)'s model assigned weights to a set of alternative numerical expressions. Surprisingly, along with weighting *one* highly (a conclusion that was supported by the empirical work), the best-fitting model assigned substantial weight to *none* as an alternative to *some*. This finding was especially surprising considering the emphasis of standard theories on scalar items that stand in entailment relationships with one another (e.g. *one* entails *some* even if it is not classically considered to be part of the scale).

In our current work, we pick up where these previous studies left off, considering the set of alternatives for implicature using the RSA model. To gain empirical traction on this issue, however, we broaden the set of scales we consider. Our inspiration for this move comes from work by van Tiel, Van Miltenburg, Zevakhina, & Geurts (2014), who examined a phenomenon that they dubbed "scalar diversity," namely the substantial difference in the strength of scalar implicature across

a variety of scalar pairs (e.g., *liked* / *loved*, *palatable* / *delicious*). Making use of this diversity allows us to investigate the ways that different alternative sets give rise to implicatures of different strengths across scales.

We begin by presenting the computational framework we use throughout the paper. We next describe a series of experiments designed to measure both the literal semantics of a set of scalar items and comprehenders’ pragmatic judgments for these same items. These experiments allow us to compare the effects of different alternative sets on our ability to model listeners’ pragmatic judgments. To preview our results: we find that standard entailment alternatives do not allow us to fit participants’ judgments, but that expanding the range of alternatives empirically (by asking participants to generate alternative messages) allows us to predict listeners’ pragmatic judgments with high accuracy.

Modeling Implicature Using RSA

We begin by giving a brief presentation of the basic RSA model. This model simulates the judgments of a pragmatic listener who wants to infer a speaker’s intended meaning m from her utterance u . For simplicity, we present a version of this model in which there is only one full recursion: The pragmatic listener reasons about a pragmatic speaker, who in turn reasons about a “literal listener.” We assume throughout that this computation takes place in a signaling game (Lewis, 1969) with a fixed set of possible meanings $m \in M$ and a fixed possible set of utterances $u \in U$, with both known to both participants. Our goal in this study is to determine which utterances fall in U .

In the standard RSA model, the pragmatic listener (denoted L_1), makes a Bayesian inference:

$$p_{L_1}(m | u) \propto p_{S_1}(u | m)p(m) \quad (1)$$

In other words, the probability of a particular meaning given an utterance is proportional to the speaker’s probability of using that particular utterance to express that meaning, weighted by a prior over meanings. This prior represents the listener’s *a priori* expectations about plausible meanings, independent of the utterance. Because our experiments take place in a context in which listeners should have very little expectation about which meanings speakers want to convey, for simplicity we assume a uniform prior where $p(m) \propto 1$.

The pragmatic speaker in turn considers the probability that a literal listener would interpret her utterance correctly:

$$p_{S_1}(u | m) \propto p_{L_0}(m | u) \quad (2)$$

where L_0 refers to a listener who only considers the truth-functional semantics of the utterance (that is, which meanings the utterance can refer to).

This model of the pragmatic speaker (denoted S_1) is consistent with a speaker who chooses words to maximize the utility of an utterance in context (Frank & Goodman, 2012), where utility is operationalized as the informativity of a particular utterance (surprisal) minus a cost:

Alternative sets				
good / excellent	liked / loved	memorable / unforgettable	palatable / delicious	some / all
excellent	loved	unforgettable	delicious	all
good	liked	memorable	palatable	most
okay	felt indifferent about	ordinary	mediocre	some
bad	disliked	bland	gross	little
horrible	hated	forgettable	disgusting	none

- Entailment items used in Experiments 2a / 3a (two-alternatives model)
- Top two alternatives added in Experiments 2b / 3b (four-alternatives model)
- Neutral item added in Experiment 2c (five-alternatives model)

Figure 1: Alternative sets for Experiments 1, 2, and 3.

$$p_{S_1}(u | m) \propto e^{-\alpha[-\log(p_{L_0}(m|u))-C(u)]} \quad (3)$$

where $C(u)$ is the cost of a particular utterance, $-\log(p_{L_0})$ represents the *surprisal* of the message for the literal listener (the information content of the utterance), and α is a parameter in a standard choice rule. If $\alpha = 0$, speakers choose randomly; as $\alpha \rightarrow \infty$, they greedily choose the highest probability alternative. In our simulations below, we treat α as a free parameter and fit it to the data.

To instantiate our signaling game with a tractable message set M , in our studies we adopt the world of restaurant reviews as our communication game. We assume that speakers and listeners are trying to communicate the number of stars in an online restaurant review (where $m \in \{1..5\}$). We then use experiments to measure three components of the model. First, to generate a set of plausible alternative messages in U , we ask participants to generate alternatives for particular scalar items (Experiment 1). Next, to measure literal semantics $p_{L_0}(m | u)$ we ask experiment participants to judge whether a message is compatible with a particular meaning (Experiment 2). Lastly, to obtain human L_1 pragmatic judgments, we ask participants to interpret a speaker’s utterances (Experiment 3).

Experiment 1: Alternative Elicitation

To examine the effect different alternative sets have on implicature, we needed a way of expanding alternative sets beyond entailment pairs. We addressed this issue by adopting a modified cloze task to elicit alternatives empirically. This design was inspired by Experiment 3 of van Tiel (2014).

Methods

Participants We recruited 30 workers on AMT. All participants were native English speakers and naive to the purpose of the experiment.

Design and procedure Participants were presented with a target scalar item from five scales, selected from van Tiel et

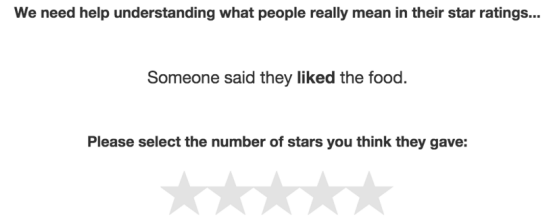


Figure 3: (Left) A trial from Experiment 2 (literal listener) with the target scalar “liked.” (Right) A trial from Experiment 3 (pragmatic listener) with the target scalar “liked.”

$z = -0.52$, $p = 0.6$). The addition of condition as a predictor did not significantly improve model fit when compared to a model without the condition variable using ANOVA ($\chi^2(2) = 0.43$, $p = 0.81$). These findings suggest that literal semantic intuitions are not effected by the set of alternatives queried in our paradigm. This finding is important because literal semantic data generated in these three conditions are used to simulate the effects of different alternative sets on implicature generation in our model.

Experiment 3: Pragmatic Listener

Experiment 3 was conducted to measure pragmatic judgments. As in Experiment 2, we include several conditions to test inferences in the presence of different alternative sets. In the two alternatives condition, participants made judgments for items included in the entailment scales. In the four alternatives condition, participants made judgments for the entailment items and also the top two alternatives elicited for each scale in Experiment 1. Including two conditions with differing alternatives allowed us to rule out the potential effects of having a larger set of alternatives during the pragmatic judgment elicitation and also provided two sets of human judgments to compare with model predictions.

Participants

We recruited 100 participants from AMT, 50 for each condition. Data for 9 participants was excluded from the two alternatives condition after participants either failed to pass two training trials or were non-native English speakers, leaving a total sample of 41 participants. In the four alternatives condition, data from 7 participants was excluded after participants either failed to pass two training trials or were not native English speakers, leaving 43 participants.

Procedure

Participants were presented with a one-sentence prompt such as “Someone said they thought the food was _____” with a target scalar item in the blank. Participants were then asked to generate a star rating representing the rating they thought the reviewer likely gave. Each participant was presented with all scalar items in a random order. The experimental setup is shown in Figure 3, right.

Results and Discussion

Figure 5 plots pragmatic listener judgment distributions for “weak” / “strong” scalar pairs (e.g. *good/excellent*). Several trends are visible. First, in each scale participants generated implicatures. They were substantially less likely to assign high star ratings to weaker scalar terms, despite the literal semantic compatibility of those terms with those states shown in Experiment 2. Second, the size of the difference between strong and weak scalar items varied considerably across scales, consistent with previous work (van Tiel et al., 2014).

To rule out the potential effects of having a larger set of alternatives during the pragmatic judgment elicitation, we fit a mixed effects model. We regressed pragmatic judgments on scale and experimental condition with subject- and word-level random effects, which was the maximal structure that converged. There were no significant differences between the two alternatives and four alternatives conditions ($\beta = 0.05$, $t(150) = 1.04$, $p = 0.3$). The addition of the condition predictor did not significantly improve model fit when compared to a model without that variable ($\chi^2(1) = 1.13$, $p = 0.29$).

Model

Using literal listener data from Experiment 2, we conducted a set of simulations with the RSA model. Each simulation kept the model constant, fitting the choice parameter α as a free parameter, but used a set of alternatives to specify the scale over which predictions were computed. We considered four different alternative sets, with empirical measurements corresponding to those shown in Figure 1: 1) the two alternatives in the classic entailment scales, 2) those two alternatives with the addition of a generic negative alternative, 3) the expanded set of four alternatives, and 4) the expanded set of five alternatives. Literal semantics for the generic negative alternative serve as a baseline “none“-style semantics in which the scalar item is only compatible with 1 star.

Model fit with human judgments was significantly improved by the inclusion of alternatives beyond the entailment items (Table 1). The two-alternatives model contained only entailment items, which, under classic accounts, should be sufficient to generate implicature, but fit to data was quite poor with these items. The addition of a generic negative element produced some gains in performance, but much higher

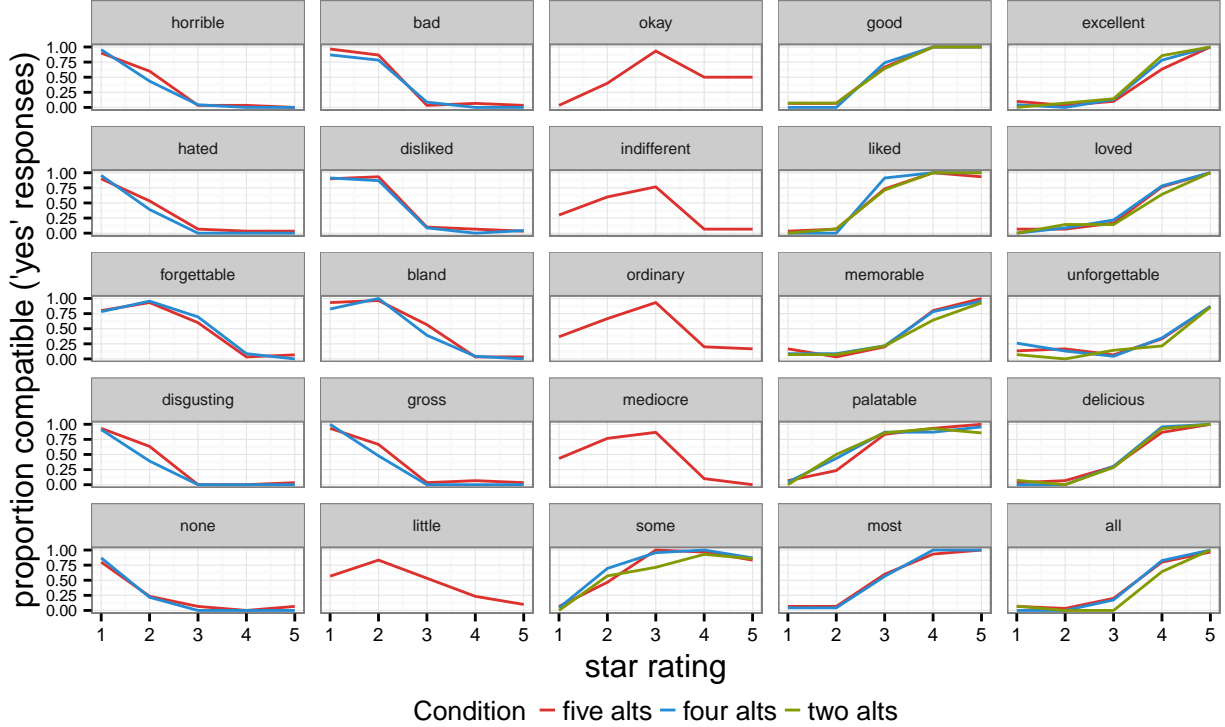


Figure 4: Literal listener judgments from Experiment 2. Proportion of participants indicating compatibility (answering “yes”) is shown on the vertical axis, with the horizontal axis showing number of stars on which the utterance was judged. Rows are grouped by scale and items within rows are ordered by valence. Colors indicate the specific condition (1a,b,c) with conditions including different numbers of items.

performance was found when we included four and five alternatives, with the alternatives derived empirically for the specific scale we used. An example fit for the five-alternatives model is shown in Figure 6.

Model	α	Two alts	Four alts
Two alts	9	0.54	0.57
Two alts + generic negative	6	0.62	0.66
Four alts	4	0.84	0.90
Five alts	4	0.86	0.90

Table 1: Model performance with fitted alpha levels. Model fit assessed through correlation with human judgments in the two conditions of Experiment 3.

General Discussion

Pragmatic inference requires reasoning about alternatives. The fundamental pragmatic computation is counterfactual: “if she had meant X, she would have said Y, but she didn’t.” Yet the nature of these alternatives has been controversial. For a few well-studied scales, a small set of logically-determined alternatives has been claimed to be all that is necessary (Horn, 1972). For other, contextually-determined inferences, the issue of alternatives has been considered relatively intractable in terms of formal inquiry (Sperber & Wilson, 1995).

In our current work, we used the rational speech act framework to investigate the set of alternatives that best allowed the model to predict pragmatic judgments across a series of different scales. We found that the best predictions came when a range of scale-dependent negative and neutral alternatives were added to the implicature computation, suggesting the importance of considering non-entailment alternatives. This work builds on previous investigations, reinforcing the claim that negative alternatives are critical for understanding implicature (Franke, 2014), and replicates and extends findings that different lexical scales produce strikingly different patterns of inference (van Tiel et al., 2014).

While improvements in model fit were substantial as we moved from two to four alternatives, we saw only a minor increase in fit from the move to five alternatives. One possible explanation is that alternatives are differentially salient in context, and in moving to larger sets we should consider weighting the alternatives differentially (as Franke, 2014 did). Preliminary simulations using weightings derived from Experiment 1 provide some support for this idea but would require further empirical work for confirmation.

The precise set of alternatives present during implicature is likely to be domain dependent. Our current empirical paradigm elicited literal semantics, pragmatic judgments, and plausible alternatives all within the restricted domain of restaurant reviews. Our measurements might have dif-

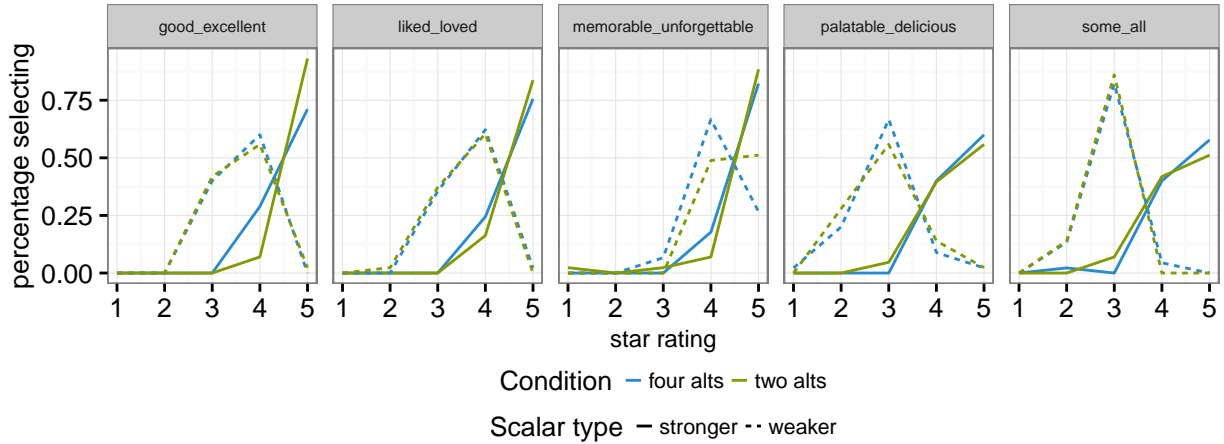


Figure 5: Pragmatic listener judgements in Experiment 3. Proportion of participants generating a star rating is shown on the vertical axis, with the horizontal axis showing number of stars on which the utterance was judged. Line type denotes condition, and colors indicate the particular scalar items. Each panel shows one scalar pair, with only entailment items (two alternatives condition) shown here for simplicity in reporting.




Figure 6: Judgments from the four-alternatives condition of Experiment 3 are plotted against model predictions using the five-alternatives data from Experiment 1. Colors show the star rating for individual judgments.

ferred substantially if we had instead grounded our ratings in a different context. Future investigations should probe the context-specificity of the weight and availability of particular alternative sets.

More broadly, considering the context- and domain-specificity of alternative sets may provide a way to unite what Grice (1975) called “generalized” (cross-context) and “particularized” (context-dependent) implicatures. Rather than being grounded in a firm distinction, we may find that these categories are simply a reflection of the effects of context on a constantly-shifting set of pragmatic alternatives.

Acknowledgements

Thanks to NSF BCS #1456077 for support, and thanks to Michael Franke, Judith Degen, and Noah Goodman for valuable discussion.

References

- Degen, J., & Tanenhaus, M. K. (2015). Processing scalar implicature: A constraint-based approach. *Cognitive Science*, 39(4), 667–710.
- Frank, M., & Goodman, N. D. (2012). Predicting pragmatic reasoning in language games. *Science*, 336(6084), 998.
- Franke, M. (2014). Typical use of quantifiers: A probabilistic speaker model. In *Proceedings of the 36th annual conference of the cognitive science society* (pp. 487–492).
- Goodman, N. D., & Stuhlmüller, A. (2013). Knowledge and implicature: Modeling language understanding as social cognition. *Topics in Cognitive Science*, 5(1), 173–184.
- Grice, H. P. (1975). Logic and conversation. In P. Cole & J. Morgan (Eds.), *Syntax and semantics* (Vol. 3). New York: Academic Press.
- Horn, L. R. (1972). *On the semantic properties of logical operators*. (PhD thesis). University of California, Los Angeles.
- Levinson, S. C. (2000). *Presumptive meanings: The theory of generalized conversational implicature*. MIT Press.
- Lewis, D. (1969). *Convention: A philosophical study*. John Wiley & Sons.
- Sperber, D., & Wilson, D. (1995). *Relevance: Communication and cognition* (2nd ed.). Oxford, UK: Blackwell.
- van Tiel, B. J. M. (2014). Quantity matters: Implicatures, typicality, and truth.
- van Tiel, B. J. M., Van Miltenburg, E., Zevakhina, N., & Geurts, B. (2014). Scalar diversity. *Journal of Semantics*, ffu017.