

Determining the alternatives for scalar implicature

Benjamin Peloquin

bpeloqui@stanford.edu

Department of Psychology

Stanford University

Michael C. Frank

mcf Frank@stanford.edu

Department of Psychology

Stanford University

Abstract

Successful communication regularly requires listeners to make pragmatic inferences - enrichments beyond the literal meaning of a speaker's utterance. For example, when interpreting a sentence such as "Bob ate some of the cookies," listeners routinely infer that Bob did not eat all of them. A Gricean account of this phenomena assumes the presence of alternatives (like "all of the cookies") with varying degrees of informativity, but it remains an open question precisely what these alternatives are. We use a computational model of pragmatic inference to test out hypotheses about how well different sets of alternatives allow us to predict scalar implicature performance across a range of different scales. Our findings suggest that human comprehenders likely consider a much broader set of alternatives beyond those entailed by the initial description.

Keywords: pragmatics; scalar implicature; bayesian modeling

Humpty Dumpty Quote...

Introduction

Successful communication requires listeners to make pragmatic inferences that go beyond the literal semantic content of speakers' utterances. For example, listeners commonly enrich the meaning of the scalar item "some" to "some but not all" in sentences like "Bob ate some of the cookies" (Grice, 1975; Horn, 1984; Levinson, 2000). These inferences, called *scalar implicatures*, have been an important test case for understanding pragmatic inferences more generally. A Gricean account of this phenomena assumes listeners reason about the meaning the speaker intended by incorporating knowledge about a) alternative scalar items a speaker could have used (such as "all") and b) the relative informativity of using such alternatives (Grice, 1975). According to this account, a listener will infer that the speaker must have intended that Bob did not eat "all" the cookies because it would have been *underinformative* for the speaker to use "some" when "all" could have been used.

But what are the alternatives that should be considered in this computation more generally? Under classic accounts of implicature, listeners consider only those words whose meaning would entail the word that is actually sent (Horn, 1972), and these alternatives enter into conventionalized or semi-conventionalized scales (Levinson, 2000). For example, because "all" entails "some," and hence is a "stronger" meaning, "all" should be considered as an alternative to "some" in implicatures. Similar scales exist for non-quantifier scales, e.g. "love" entails "like" (and hence "I liked the movie" implicates that I didn't love it).

Recent empirical evidence has called into question whether entailment scales are all that is necessary for understanding scalar implicature. For example, Degen & Tanenhaus (2015)

demonstrated that the scalar item "some" was judged less appropriate when exact numbers were seen as viable alternatives. And in a different paradigm, Tiel (2014) found converging evidence that "some" was judged to be atypical for small quantities. These data provide indirect evidence about a broader set of alternatives: since "some" is logically true of sets with one or two members, these authors argued that the presence of salient alternatives (the words "one" and "two") reduced the felicity of "some" via a pragmatic inference.

By formalizing pragmatic reasoning, computational models can help provide more direct evidence about the role that alternatives play. The "rational speech act" model (RSA) is one recent framework for understanding inferences about meaning in context (Frank & Goodman, 2012; N. D. Goodman & Stuhlmiller, 2013). RSA models frame language understanding as a special case of social cognition, in which listeners and speakers reason recursively about one another's goals. In the case of scalar implicature, a listener makes a probabilistic inference about what the speaker's most likely communicative goal was, given that she picked the quantifier "some" rather than the stronger quantifier "all." In turn, the speaker reasons about what message would best convey her intended meaning to the listener, given that he is reasoning in this way. This recursion is grounded in a "literal" listener who reasons only according to the basic truth-functional semantics of the language.

Franke (2014) used an RSA-style model to assess what alternatives a speaker would need to consider in order to produce the typicality/felicity ratings reported by Degen & Tanenhaus (2015) and Tiel (2014). In order to do this, Franke (2014)'s model assigned weights to a set of alternative numerical expressions. Surprisingly, along with weighting "one" highly (a conclusion that was supported by the empirical work), the best-fitting model assigned substantial weight to "none" as an alternative. This finding was especially surprising considering the emphasis of standard theories on scalar items that stand in entailment relationships with one another (e.g. "one" entails "some" even if it is not classically considered to be part of the scale).

In our current work, we pick up where these previous studies left off, considering the set of alternatives for implicature using the RSA model. To gain empirical traction on this issue, however, we broaden the set of scales we consider. Our inspiration for this move comes from work by Van Tiel, Van Miltenburg, Zevakhina, & Geurts (2014), who examined a phenomenon that they dubbed "scalar diversity," namely the substantial difference in the strength of scalar implicature across a variety of scalar pairs (e.g. "liked/loved," or "palat-

able/delicious.”). Making use of some of this diversity allows us to investigate the ways that different alternative sets give rise to implicatures of different strengths across scales.

We begin by presenting the computational framework we use throughout the paper. We next describe a series of experiments designed to measure both the literal semantics of a set of scalar items and comprehenders’ pragmatic judgments for these same items. These experiments allow us to compare the effects of different alternative sets on our ability to model listeners’ pragmatic judgements. To preview our results: we find that standard entailment alternatives do not allow us to fit participants’ judgements, but that expanding the range of alternatives empirically (by asking participants to generate alternative messages) allows us to model listener judgements with high accuracy.

Modeling Implicature Using RSA

We begin by giving a brief presentation of the basic RSA model. This model simulates the judgements of a pragmatic listener who wants to infer a speaker’s intended meaning m from her utterance u . For simplicity, we present a version of this model in which there is only full recursion: that is, the pragmatic listener reasons about a pragmatic speaker, who in turn reasons about a “literal listener.” We assume throughout that this computation takes place in a signaling game (Lewis, 1969) with a fixed set of possible meanings $m \in M$ and a fixed possible set of utterances $u \in U$, with both known to both participants. Our goal in this study is to determine what utterances fall in U .

In the standard RSA model, the pragmatic listener (denoted L_1), makes a Bayesian inference:

$$p_{L1}(m | u) \propto p_{S1}(u | m)p(m)$$

In other words, the probability of a particular meaning given an utterance is proportional to the speaker’s probability of using that particular utterance to express that meaning, weighted by a prior over meanings. This prior represents the listener’s *a priori* expectations about plausible meanings, independent of the utterance. Because our experiments take place in a context in which listeners should have very little expectation about which meanings speakers want to convey, for simplicity we assume a uniform prior $p(m) \propto 1$.

The pragmatic speaker in turn considers the probability that a literal listener would interpret her utterance correctly:

$$p_{S1}(u | m) \propto p_{L0}(m | u)$$

where L_0 refers to a listener who only considers the truth-functional semantics of the utterance (that is, which meanings the utterance can refer to).

This model of the pragmatic speaker is consistent with a speaker who chooses words to maximize the utility of an utterance in context (Frank & Goodman, 2012), where utility is operationalized as the informativity of a particular utterance (surprisal) minus a cost:

Alternative sets				
good / excellent	liked / loved	memorable / unforgettable	palatable / delicious	some / all
excellent	loved	unforgettable	delicious	all
good	liked	memorable	palatable	most
okay	felt indifferent about	ordinary	mediocre	some
bad	disliked	bland	gross	little
horrible	hated	foregettable	disgusting	none

Entailment items used in Experiments 1a,b
 Top two alternatives added in Experiments 3a,b
 Neutral item added in Experiment 4

Figure 1: Stimuli for Experiments 1, 2, and 3

$$p(u | m) \propto e^{-\alpha(-\log(p_{L0}(m|u)) - C(u))},$$

where $C(u)$ is the cost of a particular utterance, $-\log(p_{L0})$ represents the *surprisal* for the literal listener (the information content of the utterance), and α is a parameter in a standard choice rule. If $\alpha = 0$, speakers choose randomly and as $\alpha \rightarrow \infty$, they greedily choose the highest probability alternative. In our simulations below, we treat α as a free parameter and fit it to the data.

To instantiate our signaling game with a tractable message set M , in our studies we adopt a food-review paradigm: we assume that speakers and listeners are trying to communicate the number of stars in an online restaurant review (where $m \in \{1, 2, 3, 4, 5\}$). We then use experiments to measure three components of the model. First, to measure literal semantics $p_{L0}(m | u)$ (we ask experiment participants to judge whether a message is applicable to a particular meaning (Experiment 1). Second, to generate a set of plausible alternative messages in U , we elicit alternatives directly (Experiment 2). Lastly, to obtain human L_1 pragmatic judgments, we ask participants to interpret a speaker’s utterances (Experiment 3).

Experiment 1: Literal listener task

Experiment 1 was conducted to approximate literal listener semantic distributions $p_{L0}(m | u)$ for five pairs of scalar items taken from Tiel (2014). We include three conditions in Experiment 1, corresponding to the sets of alternatives within a scale that participants were presented with: two alternatives (“entailment”), four alternatives, and five alternatives. The two-alternative entailment condition makes a test of the hypothesis that the two members of the classic Horn scale (Horn, 1972) are the only alternatives necessary to predict the strength of listeners’ pragmatic inference. The four- and five-alternative conditions then add successively more alternatives to test whether including a larger number of alternatives will

We're trying to better understand what people really mean in their star ratings

When rating the food someone gave this many stars:



Do you think the person **liked** the food?

No ☐ Yes ☐

We need help understanding what people really mean in their star ratings...

Someone said they **liked** the food.

Please select the number of stars you think they gave:




Figure 2: (Left) A trial from Experiment 1 (literal listener) with the target scalar ‘liked.’ (Right) A trial from Experiment 3 (pragmatic listener) with the target scalar ‘liked.’

increase model fit.¹

Methods

Participants In each condition we recruited 30 participants from Amazon Mechanical Turk (AMT). In the two-alternative condition, two participants were excluded because they reported having a native language other than English for a final sample of 28 participants. In the four-alternative condition, six participants were excluded for either failing to pass two training trials a non-English native language, leaving a total sample of 24 participants. In the five-alternative condition, there were no data exclusions due to training failures or native language requirements, leaving a total sample of 30 participants.

Design and procedure Figure 2, left, shows the experimental setup. Participants were presented with a target scalar item and a star rating (1–5 stars) and asked to judge the compatibility of the scalar item and star rating. Compatibility was assessed through a binary “yes/no” response to a question of the form “Do you think that the person thought the food was ____?” where a target scalar was presented in the blank. Each participant saw all scalar item and star rating combinations for their particular condition (see Figure 1), in a random order.

The two-alternatives condition included only the scalar pairs from Tiel (2014). The four-alternatives condition included the two scalar items plus the top two alternatives generated for each scalar family by participants in Experiment 2. The five-alternatives condition included the four previous items plus one more “neutral” item chosen from those alternatives generated in Experiment 2. See Figure 1 for the complete list of alternatives used in each condition.

Results and Discussion

Figure 2 plots literal listener $p_{LO}(m|u)$ distributions for the three conditions. Several trends are visible. First, in each scale, the alternatives spanned the star scale, such that there were alternatives that were highly compatible with both the lowest and highest numbers of stars. Second, there was clear

variability between scalar families. For example, compatibility judgments for the top items in the “memorable / unforgettable” scale were more similar than those for “good / excellent” or “liked / loved.” Finally, there was substantial consistency in ratings for items that were repeated across experiments, suggesting that this paradigm elicited stable judgments from participants.

Experiment 2: Alternative Elicitation

To elicit empirical alternatives for the scales we used in Experiment 1, we adopted a modified cloze task inspired by Experiment 3 of Tiel (2014).

Methods

Participants We recruited 30 workers on AMT. All participants were native English speakers and naive to the purpose of the experiment.

Design and procedure Participants were presented a target scalar term from our original entailment set embedded in a sentence such as, “In a recent restaurant review someone said they thought they the food was ____,” with a target scalar presented in the blank. Participants were then asked to generate plausible alternatives by responding to the question, “If they’d felt differently about the food, what other words could they have used instead of ____?” They were prompted to generate three unique alternatives.

Results and Discussion

Figure ?? an example alternative set for the scalar items “liked” and “loved” (combined). Alternative distributions for the other scalar pairs (e.g., “good/excellent”, “memorable/unforgettable”) were similarly long-tailed.

Experiment 3: Pragmatic Listener

In Experiment 3, we asked participants to make pragmatic listener judgements.

CONDITIONS

Participants

We recruited 150 participants from AMT, 50 for each condition. Data for 7 participants in the two-alternative condition were excluded after participants either failed to pass two

¹Note that alternatives in the four- and five-alternatives conditions were chosen on the basis of Experiment 2, which was run chronologically after the two-alternative condition; all literal listener experiments are grouped together for simplicity in reporting.

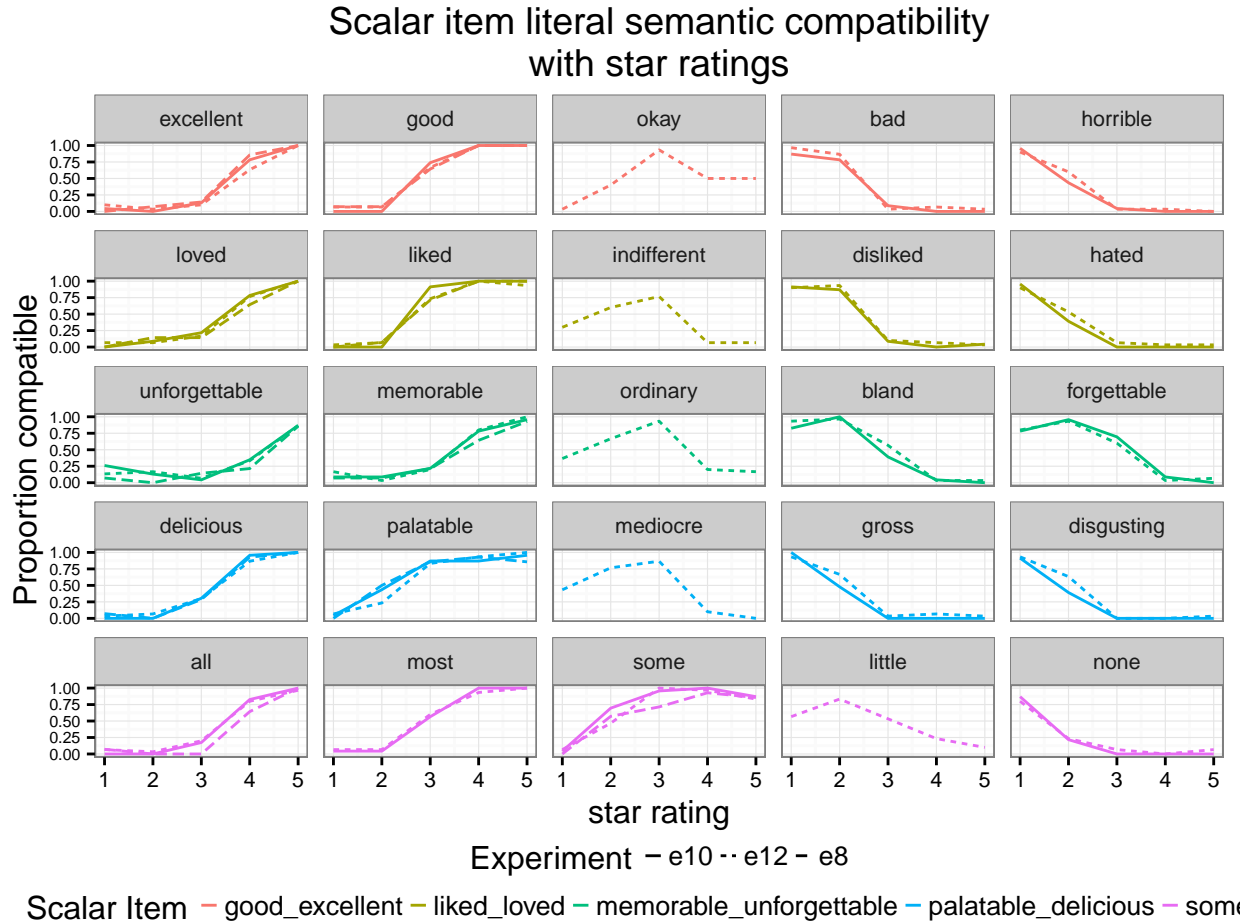


Figure 3: Literal listener judgements from Experiments 1a, 3a, and 4. Proportion of participants indicating acceptability is shown on the vertical axis, with the horizontal axis showing number of stars on which the utterance was judged. Each line-type shows a different experiment, and colors indicate the different items in the scale (with the experiments including different numbers of items). Each panel shows one scalar pair.

training trials or reported a non-English native language, leaving a total sample of 43 participants. In the four-alternative condition, data from five participants was thrown out after participants either failed to pass two training trials or were not native English speakers, leaving a total sample of 45 participants.

Procedure

Participants were presented with a one-sentence prompt containing a target scalar item such as “Someone said they thought the food was _____.” Participants were then asked to generate a star rating representing the rating they thought the reviewer likely gave. Each participant was presented with all scalar items in a random order. The experimental setup is shown in Figure ??, right.

Results and Discussion

Model Results

Using literal semantic data from Experiment 1 we conducted three simulations with our model. Each simulation used the

specific literal semantic data to specify the scale representation available to our model. The “entailment” model used only the original pair of scalar items pulled from van Tiel (2014). These data were measured in Experiment 1a. The “Top two” model extended the original set of alternatives with the top two alternatives generated in Experiment 2. These data were measured in Experiment 3a. The “Full” model included the full set of alternatives, including a neutral valenced alternative. These data were measured in Experiment 4.

General Discussion

By varying the type of scale representations available to our Bayesian model we investigated the effects of alternatives on scalar implicature. Model fit with human judgement was significantly improved by the inclusion of alternatives beyond the typical “strong” and “weak” scalar items. In fact, we found that both neutral and negative valence scalar items contribute to human-like implicature generation within our framework.




Figure 5: Pragmatic listener judgements for our entailment scalar items. Proportion of participants generating a star rating is shown on the vertical axis, with the horizontal axis showing number of stars on which the utterance was judged. Each line-type shows a different experiment, and colors indicate the stronger (ie excellent) or weaker (ie good) scalar items for each scale. Each panel shows one scalar pair.




Figure 4: Caption goes here

Acknowledgements

Thanks to NSF BCS XYZ. Thanks to Michael Franke, Judith Degen, and Noah Goodman.

References

Degen, J., & Tanenhaus, M. K. (2015). Processing scalar implicature: A constraint-based approach. *Cognitive Science*, 39(4), 667–710.

Frank, M., & Goodman, N. (2012). Predicting pragmatic reasoning in language games. *Science*, 336(6084), 998.

Franke, M. (2014). Typical use of quantifiers: A probabilistic

speaker model. In *Proceedings of the 36th annual conference of the cognitive science society* (pp. 487–492).

Goodman, N. D., & Stuhlmiller, A. (2013). Knowledge and implicature: Modeling language understanding as social cognition. *Topics in Cognitive Science*, 5(1), 173–184.

Grice, H. P. (1975). Logic and conversation. In P. Cole & J. Morgan (Eds.), *Syntax and semantics* (Vol. 3). New York: Academic Press.

Horn, L. R. (1972). *On the semantic properties of logical operators*. (PhD thesis). University of California, Los Angeles.

Horn, L. R. (1984). Toward a new taxonomy for pragmatic inference: Q-based and R-based implicature. *Meaning, Form, and Use in Context*, 42.

Levinson, S. C. (2000). *Presumptive meanings: The theory of generalized conversational implicature*. MIT Press.

Lewis, D. (1969). *Convention: A philosophical study*. John Wiley & Sons.

Tiel, B. van. (2014). Quantity matters: Implicatures, typicality, and truth.

Van Tiel, B., Van Miltenburg, E., Zevakhina, N., & Geurts, B. (2014). Scalar diversity. *Journal of Semantics*, ffu017.

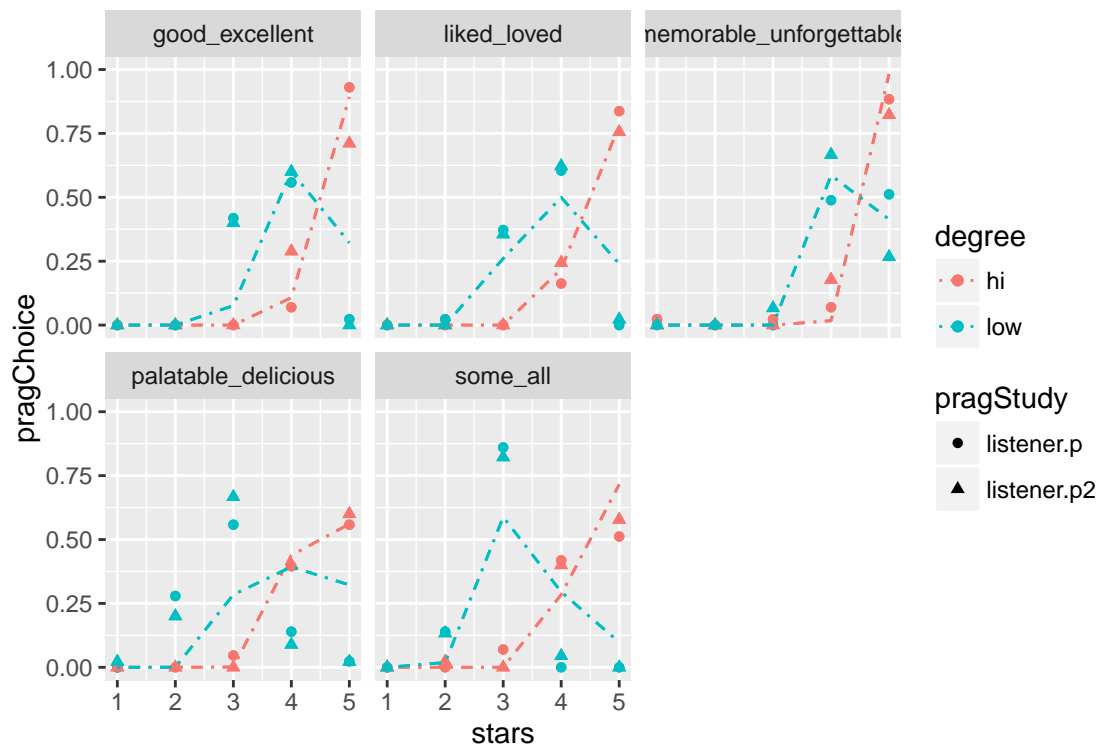


Figure 6: The left panel shows improved model fit as scale representations are enriched with more scalar items. Correlations are computed using pragmatic judgment data from Experiments 1b and 3b. The right panel plots model predictions using full symmetric scales versus human judgments from Experiment 3b.